

# 5660 HW1 Report

## LangChain Receipt QA Agent

Name: Liu Sixiao  
Student ID:1155247019

### 1. Overview

This project builds a LangChain-based system that takes multiple supermarket receipt inputs (extracted text through OCR) with user queries, and returns answers for two supported queries:

- **Query 1:** How much money did I spend in total for these bills?
- **Query 2:** How much would I have had to pay without the discount?
- **Irrelevant query handling:** refuse to answer if the input question is unrelated to the receipts

### 2. System Design

#### Input representation

All receipts are concatenated into a single string *receipt\_content*. To avoid cross-receipt contamination, I enforce explicit session boundaries:

----begin---- ... receipt text ... ----end----

This boundary design makes it easier for the model to treat each receipt as an independent session and reliably extract totals/discounts per session.

#### LLM and LangChain stack

- LLM: ChatGoogleGenerativeAI (Gemini), with **temperature=0** for stable outputs.
- Prompting: *ChatPromptTemplate.from\_messages()* using a system role + user instruction.
- Composition: LCEL pipeline pattern *prompt | llm* for clean modular chaining.
- Output handling: an extra normalization step to guarantee **float-parseable** numeric outputs.

### 3. Chains and Prompt Engineering

I implemented one dedicated pipeline for each query, plus a shared idea of output normalization to ensure the final answer can be directly cast to *float*.

#### 3.1 Query 1: Total Amount Spent (Summation)

Chain definition: **full\_query1\_pipeline = prompt1 | llm**.

Prompt1 instructs the model to scan each receipt session (bounded by begin/end), extract the session total (e.g., TOTAL/AMOUNT PAID), and sum across sessions.

The output must contain **only the final numeric result** with no reasoning text.

#### Output normalization:

To prevent formats like “Total: \$123.45”, I apply a second prompt (result-transformation) that outputs only a numeric string.

### 3.2 Query 2: Payment Without Discount (Avoid Double Counting)

Query 2 is harder because discounts can appear in multiple places (item lines, discount sections, or subtotals). I use a two-step approach: (A) discount-focused extraction to reduce noise and then (B) calculation prompt that adds back discounts carefully without double-counting.

- **Step A (filter):** Extract only discount-related lines/sections plus each session total.
- **Step B (compute):** Use prompt2 to compute “pay without discount” while checking for discount duplication.
- **Normalization:** Apply a final transformer prompt to output a float-compatible numeric string.

## 4. Query Handling and Rejection

To satisfy the requirement of rejecting irrelevant queries, I implemented a **prompt-level rejection policy** inside the system messages of both Query 1 and Query 2 pipelines. Conceptually, the model first checks whether the input is receipt-related; if not, it should return a refusal response instead of fabricating an answer.

In the notebook, I also add a simple sanity test by sending an unrelated question to the pipeline (e.g., “What’s the weather today in HK?”)

```
Irrelevant_response = full_query2_pipeline.invoke("What's the weather today in HK?")
```

## 5. Evaluation

I validate outputs using a lightweight test function that compares the model answer against ground-truth totals with a tolerance of  $\pm 2$  to account for OCR noise and rounding differences.

This ensures: (1) outputs are numeric and parseable, and (2) totals are close to the expected sums.

## 6. Discussion and Future Improvements

- **Strengths:** receipt session boundaries, task-specific prompts, and strict numeric normalization improve stability.
- **Discount realism:** in real receipts, discounts often appear as negative lines (e.g., -1.50). A future improvement is to classify negative amounts using keywords (discount,refund,return).
- **Consistency:** unify delimiter strings across all steps to avoid ambiguity.
- **Agent completion:** wrap Query1/Query2 as LangChain tools and add a router to fully satisfy query rejection.

## 7. Conclusion

Overall, the solution demonstrates a clean post-OCR LangChain design: multi-receipt text structuring, two task pipelines for Query1 and Query2, output normalization for float parsing, and automatically rejecting irrelevant queries.