

Moltbook Social Agent Report

Name: Liu Sixiao
Student ID:1155247019

Overview

This report documents the design, autonomy logic, and Moltbook interaction evidence for an agentic AI system that authenticates with Moltbook, navigates the /m/ftec5660 submolt, and performs social actions (subscribe, upvote, and comment) via the public REST API.

1. Problem Requirements

The agent must:

- (i) authenticate with the Moltbook API using an API key
- (ii) subscribe to the course submolt /m/ftec5660
- (iii) upvote and comment on the pinned welcome post .

2. System Overview

The system is organized into three layers:

- (1) Tool layer: thin wrappers over Moltbook REST endpoints (feed, search, list submolts, get posts from a submolt, subscribe, create post/comment, and upvote).
- (2) Agent layer: an LLM configured for tool-calling; it receives a task instruction and decides which tools to call, with arguments, to complete the task.
- (3) Runtime loop: a controller that executes tool calls append too output back into the conversation history, and repeats until the LLM returns a final message with no further tool calls.

3. Tooling Layer

Each tool is a LangChain @tool function that calls a single REST endpoint using requests with an Authorization Bearer token. The most important tools for the required tasks are: get_submolt_posts (to inspect /m/ftec5660 content), subscribe (to subscribe the submolt), upvote_post unsubscribe and comment_post.

```
@tool
def get_submolt_posts(name: str, sort: str="new") -> dict:
    """get posts from a submolt"""
    r = requests.get(
        f"{BASE_URL}/posts",
        headers=HEADERS,
        params={"submolt": name, "sort": sort},
        timeout=25
    )
    return r.json()
```

```
@tool
def subscribe(name: str) -> dict:
    """Subscribe a submolt."""
    r = requests.post(
        f"{BASE_URL}/submolts/{name}/subscribe",
        headers=HEADERS,
        timeout=15
    )
    return r.json()
```

```
@tool
def unsubscribe(name: str) -> dict:
    """Subscribe a submolt."""
    r = requests.delete(
        f"{BASE_URL}/submolts/{name}/subscribe",
        headers=HEADERS,
        timeout=15
    )
    return r.json()
```

4. Agent Design and Architecture

The agent uses a tool-calling LLM and a deterministic controller loop. The controller maintains a history of messages (user task, tool outputs, and model messages). On each turn, the LLM is invoked with the current history; if the model responds with tool_calls, the controller executes them synchronously, logs inputs/outputs, and appends the results back into history. This continues until the model returns a final natural-language response without tool calls.

This design separates (a) decision-making (LLM) from (b) safe execution (controller + tool wrappers), enabling observable, replayable interactions.

Core runtime loop excerpt:

```
# =====
# Main agent loop
# =====
for turn in range(1, max_turns + 1):
    log("TURN", f"Turn {turn}/{max_turns} started")
    turn_start = time.time()

    response = agent.invoke(history)
    history.append(response)

    if verbose:
        log("LLM", "Model responded")
        log("LLM.CONTENT", response.content or "<empty>")
        log("LLM.TOOL_CALLS", pretty(response.tool_calls or []))

    # =====
    # STOP CONDITION
    # =====
    if not response.tool_calls:
        elapsed = round(time.time() - turn_start, 2)
        log("STOP", f"No tool calls – final answer produced in {elapsed}s")
        return response.content
# ... (truncated)
```

5. Decision Logic and Autonomy Level

Autonomy is task-driven and bounded by the available tool set. For each task, the LLM chooses an action sequence based on tool affordances and returns structured tool calls. For the homework's required workflow, the intended plan is:

- Identify the target submolt (ftec5660) and fetch its posts (get_submolt_posts).
- Subscribe to /m/ftec5660 (subscribe).
- Locate the target post id (either directly provided by the homework or inferred as the top/pinned post from sorted results).
- Upvote the post (upvote_post).
- Add a short course-relevant comment (comment_post).

The controller executes exactly the actions the model requests, but only through pre-defined tools. This provides a practical balance: the model is autonomous in choosing steps and arguments, while execution is constrained to vetted endpoints.

6. Interaction Evidence (Logs)


The notebook includes verbose run logs that show tool calls, arguments, and returned JSON. Below is an excerpt where the agent asks for the most upvoted post in /m/ftec5660 and correctly retrieves the welcome post id used by the homework.

```
"sort": "top",
"name": "ftec5660"
}
[11:12:07] [TOOL.RESULT] get_submolt_posts finished (success) in 9.37s
[11:12:07] [TOOL.OUTPUT] {
  "success": true,
  "posts": [
    {
      "id": "47ff50f3-8255-4dee-87f4-2c3637c7351c",
      "title": "Welcome to FTEC5660 ",
      "content": "Use this submolt to share questions, notes, experiments, and insights related to the",
      "type": "text",
      "author_id": "f8a80401-bdff-4c0d-bc92-076af920cc2f",
      "author": {
        "id": "f8a80401-bdff-4c0d-bc92-076af920cc2f",
        "name": "BaoNguyen",
        "description": "Baos agent",
        "avatarUrl": null,
        "karma": 33,
        "followerCount": 1,
        "followingCount": 1,
```

The system also demonstrates resilience to transient failures via repeated attempts: in the following excerpt, a timeout occurs, then the agent retries on the next turn.

```
"args": {
  "name": "ftec5660"
},
"id": "ccf6158c-eb37-4d76-9536-184b6450ff47",
"type": "tool_call"
}
]
[12:14:33] [TOOL] [1] Calling `get_submolt_posts`
[12:14:33] [TOOL.ARGS] {
  "name": "ftec5660"
}
[12:14:58] [TOOL.RESULT] get_submolt_posts finished (error) in 25.15s
[12:14:58] [TOOL.OUTPUT] {
  "error": "HTTPSPool(host='www.moltbook.com', port=443): Read timed out. (read timeout=25)"
}
```

Here is the screenshots of Moltbook interaction, shows two comments posted by my agent



7. Reliability Considerations and Improvements

Observed issues include occasional request timeouts and internal server-side 500 errors.

Practical improvements are:

- (i) Add `r.raise_for_status()` and structured error turns in every tool
- (ii) Implement exponential back off on 429/timeouts
- (iii) Cap tool output size before passing back to the LLM(e.g.,keep only top-K posts and essential fields);
- (iv) Add a small planner step that pins the required post id and verifies success (e.g., check vote count or read-back).

8. Conclusion

This agent implements an end-to-end agentic workflow on a real social platform for agent: authenticated API access, submolt navigation, and social actions. The architecture cleanly separates tools, decision logic, and execution control, and the logs provide auditable evidence of correct API usage, including discovery of the required welcome post id in `/m/fttc5660`.